

Load the dataset

In [1]:

```
from sklearn.datasets import load_boston
boston = load_boston()
```

Dataset characteristics

In [2]:

```
print(boston.DESCR)
```

```
.. _boston_dataset:
```

```
Boston house prices dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
    :Number of Instances: 506
```

```
    :Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.
```

```
    :Attribute Information (in order):
```

```
        - CRIM      per capita crime rate by town
```

```
        - ZN        proportion of residential land zoned  
d for lots over 25,000 sq.ft.
```

```
        - INDUS     proportion of non-retail business a  
cres per town
```

```
        - CHAS      Charles River dummy variable (= 1 i  
f tract bounds river; 0 otherwise)
```

```
        - NOX       nitric oxides concentration (parts  
per 10 million)
```

```
        - RM        average number of rooms per dwellin  
g
```

```
        - AGE       proportion of owner-occupied units  
built prior to 1940
```

- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/> (<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>)

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.

- Quinlan,R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [3]:

```
print(f'  Data shape = {boston.data.shape}')  
print(f'Target shape = {boston.target.shape}')
```

```
  Data shape = (506, 13)  
Target shape = (506,)
```

In [4]:

```
boston.feature_names
```

Out[4]:

```
array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',  
      'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

Load data into a dataframe and explore

In [5]:

```
import pandas as pd

boston_df = pd.DataFrame(boston.data,
                          columns=boston.feature_names)

boston_df.head()
```

Out[5]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	F
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

target value MEDV is missing, we create a new column for MEDV and add into data frame

In [6]:

```
boston_df['MEDV'] = pd.Series(boston.target)

boston_df.head()
```

Out[6]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	F
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

In [7]:

```
boston_df.describe()
```

Out[7]:

	CRIM	ZN	INDUS	CHAS	NOX	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284854
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.707115
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561542
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.884612
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.209519
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.629166
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.787734


Data visualizations

We use heatmap to visualize the linear relationship between variables. In this case, we are only looking at the realtionship to MEDV column.

In [8]:

```
sample_df = boston_df.sample(frac=0.4) # sample data to 40%
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

plt.figure(figsize=(15,8))
correlation_matrix = sample_df.corr().round(2)
# annot = True to print the values inside the square
axe=sns.heatmap(data=correlation_matrix, annot=True,cmap="Reds")
axe.set_title('Boston Housing Price Heatmap');
```





From heatmap, we can see: CHAS, DIS, B have weak correlation between MEDV with coefficient of 0.16, 0.28 and 0.4 respectively.

In [9]:

```
sns.set(font_scale=2)
sns.set_style('whitegrid')
```

In [18]:

```
# We used the full data set to plot these graphs because there is no  
for feature in boston.feature_names:  
    plt.figure(figsize=(16, 4))  
    sns.scatterplot(data=boston_df, x=feature, y='MEDV',  
                    hue='MEDV',  
                    palette='cool', legend=False)
```



Split the data for training and testing

In [16]:

```
from sklearn.model_selection import train_test_split
import numpy as np
boston_data_new = np.delete(boston.data,3,1)
# We found that CHAS contains only variables 0 and 1, which may hurt
X_train, X_test, y_train, y_test = \
    train_test_split(boston_data_new, boston.target, random_state=5)

print(f'Shape of training set = {X_train.shape}')
print(f'Shape of testing set = {X_test.shape}')
```

Shape of training set = (379, 12)

Shape of testing set = (127, 12)

Train the model

In [13]:

```
from sklearn.linear_model import LinearRegression

linear_regression = LinearRegression()
linear_regression.fit(X=X_train, y=y_train)

for i, name in enumerate(np.delete(boston.feature_names, 3, 0)):
    print(f'{name:>11}: {linear_regression.coef_[i]:24.20f}')

print()
print(f'y-intercept: {linear_regression.intercept_:23.20f}')
```

```
CRIM:  -0.15913839634194795747
ZN:    0.04890147232976368302
INDUS: 0.01445084249427397612
NOX:   -13.69906809612140463628
RM:    3.95110016068859470906
AGE:   -0.00539332527178060499
DIS:   -1.45933358556167180886
RAD:    0.36650893681391799594
TAX:   -0.01395246623283373857
PTRATIO: -0.95282684633969105814
B:      0.01359438499293804704
LSTAT: -0.53257535640300879276
```

```
y-intercept: 32.62557602342805296303
```

Test the model

In [14]:

```
import math
from sklearn import metrics

predicted = linear_regression.predict(X_test)
expected  = y_test

r2 = metrics.r2_score(expected, predicted)
r  = math.sqrt(r2)

print(f'coefficient of determination = {r2:.1f}')
print(f'      correlation coefficient = {r:.1f}')
```

```
coefficient of determination = 0.7
      correlation coefficient = 0.8
```

Visualize the expected vs. predicted prices

```
In [15]:
```

```
df = pd.DataFrame()

df['Expected'] = pd.Series(expected)
df['Predicted'] = pd.Series(predicted)


figure = plt.figure(figsize=(15, 10))

axes = sns.scatterplot(data=df, x='Expected', y='Predicted',
                      hue='Predicted', palette='cool',
                      legend=False)

start = min(expected.min(), predicted.min())
end = max(expected.max(), predicted.max())

axes.set_xlim(start, end)
axes.set_ylim(start, end)

line = plt.plot([start, end], [start, end], 'k--')
```



Load the dataset

In [1]:

Dataset characteristics

In [2]:

```
.. _boston_dataset:
```

Boston house prices dataset

****Data Set Characteristics:****

 :Number of Instances: 506

 :Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

 :Attribute Information (in order):

 - CRIM per capita crime rate by town

 - ZN proportion of residential land zone

d for lots over 25,000 sq.ft.

 - INDUS proportion of non-retail business a

res per town

 - LSTAT lower status of the population

In [3]:

```
Data shape = (506, 13)
```

```
Target shape = (506,)
```

In [4]:

Out[4]:

```
array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',  
      'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')
```

Load data into a dataframe and explore

In [5]:

Out[5]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	F
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

target value MEDV is missing, we create a new column for MEDV and add into data frame

In [6]:

Out[6]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	F
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

In [7]:


Out[7]:

	CRIM	ZN	INDUS	CHAS	NOX	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.281659
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.707115
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561541
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.884610
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208509
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.629166
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.782705

Data visualizations

We use heatmap to visualize the linear relationship between variables. In this case, we are only looking at the relationship to MEDV column.

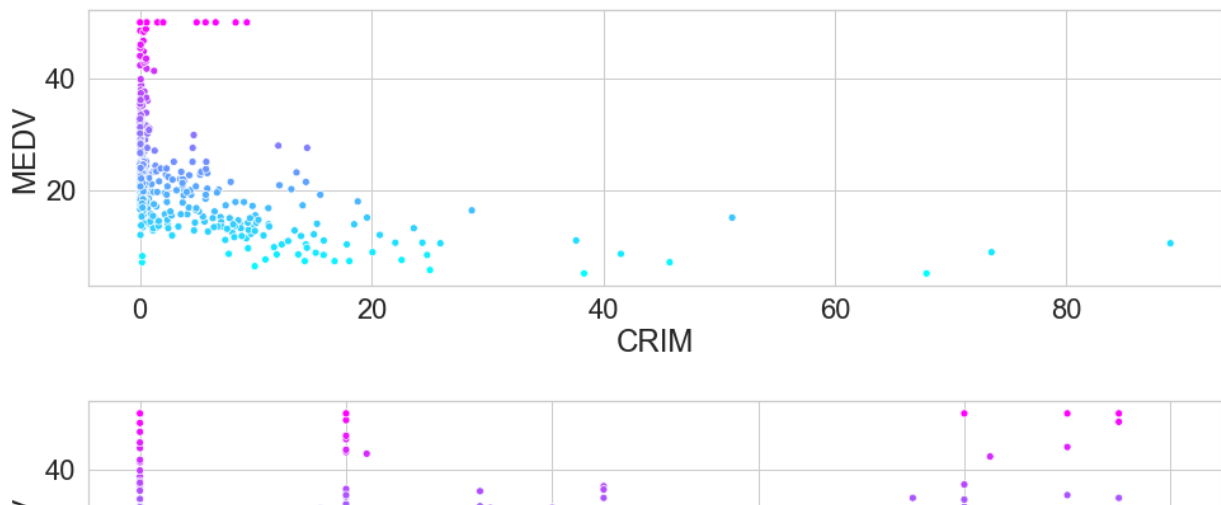
In [8]:














From heatmap, we can see: CHAS, DIS,B have weak correlation between MEDV with coefficient of 0.16, 0.28 and 0.4 respectively.

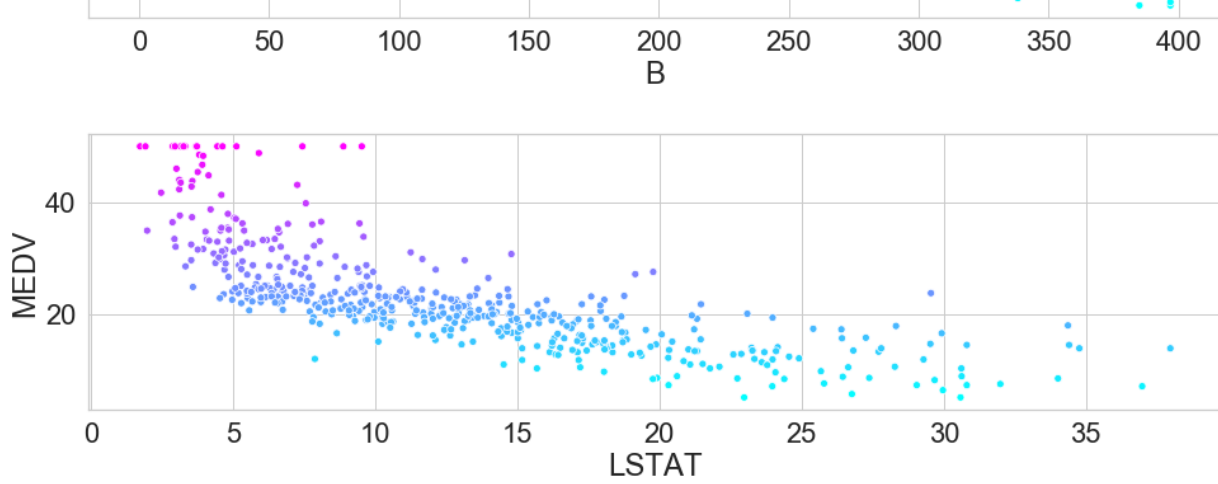
In [9]:

In [18]:









Split the data for training and testing

In [16]:

```
Shape of training set = (379, 12)
Shape of testing set = (127, 12)
```

Train the model

In [13]:

```
CRIM:  -0.15913839634194795747
ZN:    0.04890147232976368302
INDUS:  0.01445084249427397612
NOX:   -13.69906809612140463628
RM:     3.95110016068859470906
AGE:   -0.00539332527178060499
DIS:   -1.45933358556167180886
RAD:    0.36650893681391799594
TAX:   -0.01395246623283373857
PTRATIO: -0.95282684633969105814
B:      0.01359438499293804704
LSTAT: -0.53257535640300879276
```

```
y-intercept: 32.62557602342805296303
```

Test the model

In [14]:

```
coefficient of determination = 0.7
correlation coefficient = 0.8
```

Visualize the expected vs. predicted prices

In [15]:

